

Empresas dirigidas por datos: Amazon.com

Amazon killed Borders and RadioShack

amazon.com



Empresas dirigidas por datos: Netflix.com

Netflix killed Blockbuster




Recomendaciones?



Recomendación Vs Búsqueda



consulta



respuesta



Predicción de la toma de decisiones




←
sugerencia
personalizada

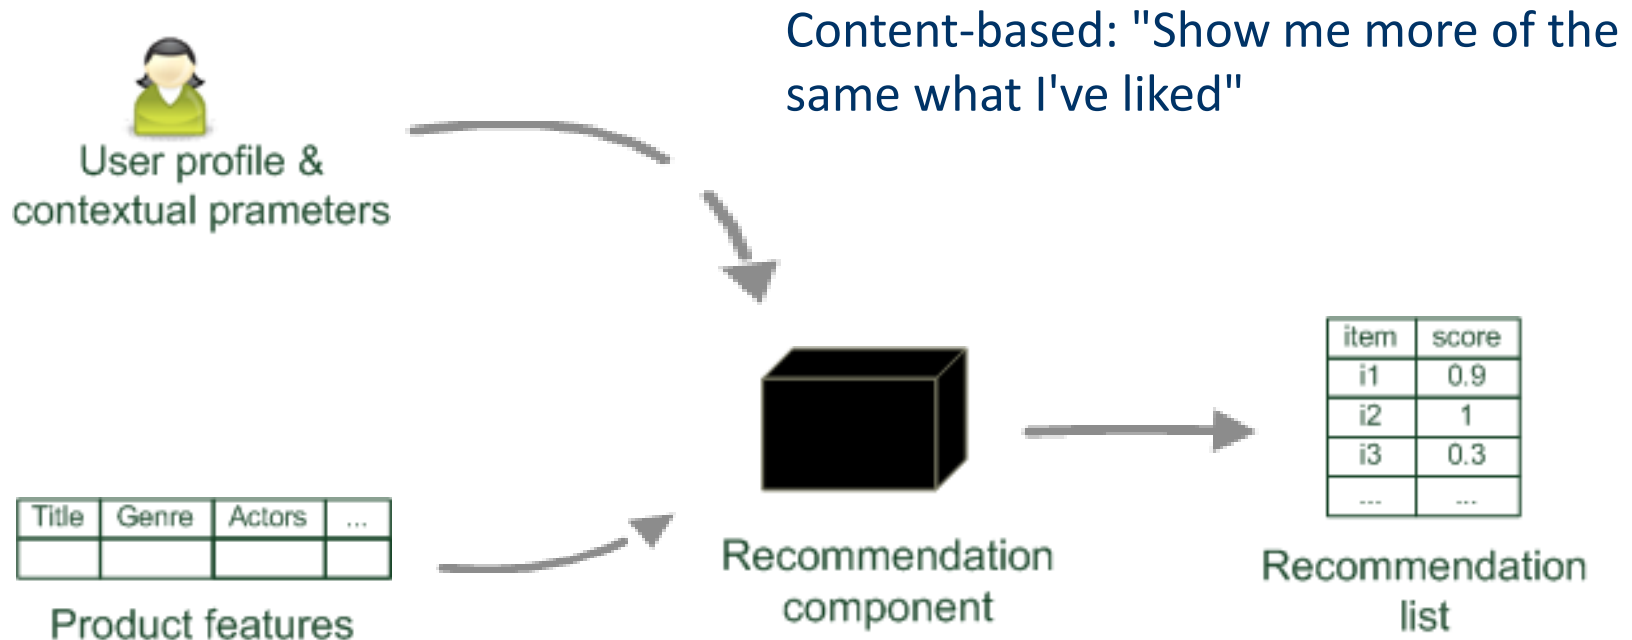



El problema de la Recomendación

Problema: Predecir la utilidad de un item
Para un usuario!!!

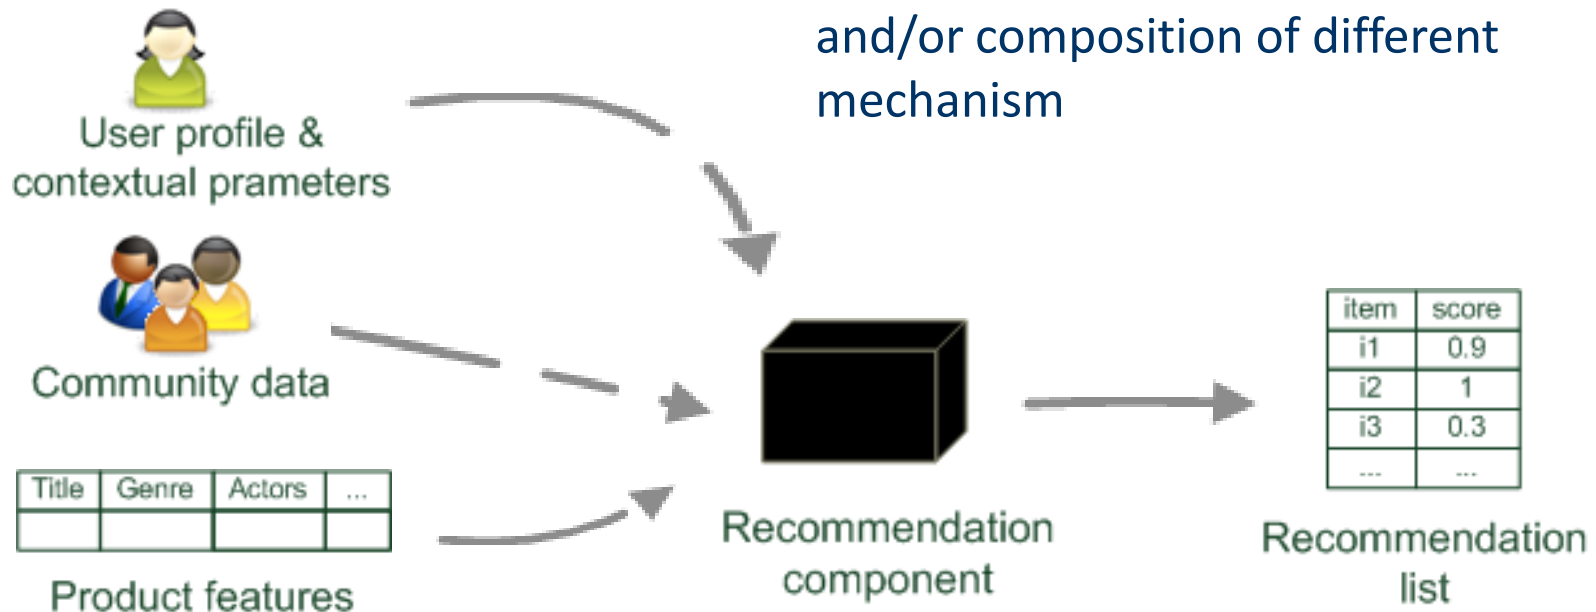


Fuente: Recommender Systems: An introduction
(Cambridge University Press)





Hybrid: combinations of various inputs
and/or composition of different
mechanism



Ejercicio: ¿Cómo genera amazon.com sus recomendaciones?

amazon.com

amazon.com

Recommended for You

Amazon.com has new recommendations for you based on [items](#) you purchased or told us you own.



[The Little Big Things: 163 Ways to Pursue EXCELLENCE](#)



[Fascinate: Your 7 Triggers to Persuasion and Captivation](#)

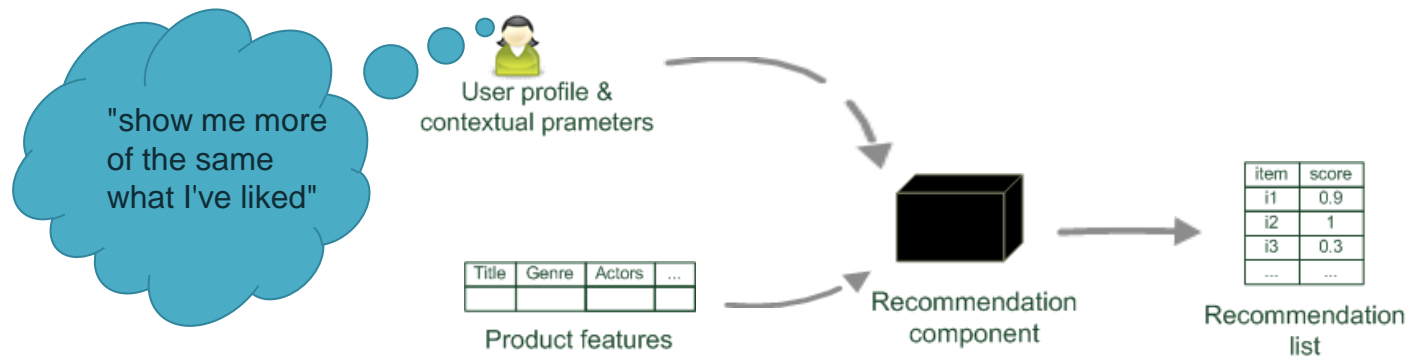


[Sherlock Holmes \[Blu-ray\]](#)



[Alice in Wonderland \[Blu-ray\]](#)

- What do we need:
 - some information about the available items such as the genre ("content")
 - some sort of *user profile* describing what the user likes (the preferences)
- The task:
 - learn user preferences
 - locate/recommend items that are "similar" to the user preferences



¡ Score = Utilidad !

► ¿Cómo construimos la función de utilidad?:

1. Identificar los atributos relevantes de las alternativas
2. Identificar los valores de los atributos anteriores
3. Estimar las preferencias del individuo sobre los valores de los atributos.
4. Matemáticamente: Función lineal sobre los valores de los atributos. Dado individuo c y alternativa a :

$$u(c, a) = \sum_k \beta_{c,k} x_{a,k}$$

Con K indicando el cto de valores de todos los atributos de a .

Conjunto de alternativas

Instancias

Atributos

Producto	Autor	Fecha	Precio	Género	Utilidad
Libro1	Valor	Valor	Valor	Ficción	10
Libro2	Valor	Valor	Valor	Histórico	3
Libro3	Valor	Valor	Valor	Novela	7
LibroN	Valor	Valor	Valor	Novela	1



Juicio de
utilidad

Técnica: Conjoint Analysis

Conjunto de datos (dataset)

Instancias

Atributos

Producto	Autor	Fecha	Precio	Género	Interacción
Libro1	Valor	Valor	Valor	Ficción	Comprado
Libro2	Valor	Valor	Valor	Histórico	Visitado
....
LibroN	Valor	Valor	Valor	Novela	Comentado

Inferencia de la Utilidad

Producto	Autor	Fecha	Precio	Género	Utilidad
Libro1	Valor	Valor	Valor	Ficción	10
Libro2	Valor	Valor	Valor	Histórico	5
LibroN	Valor	Valor	Valor	Novela	7

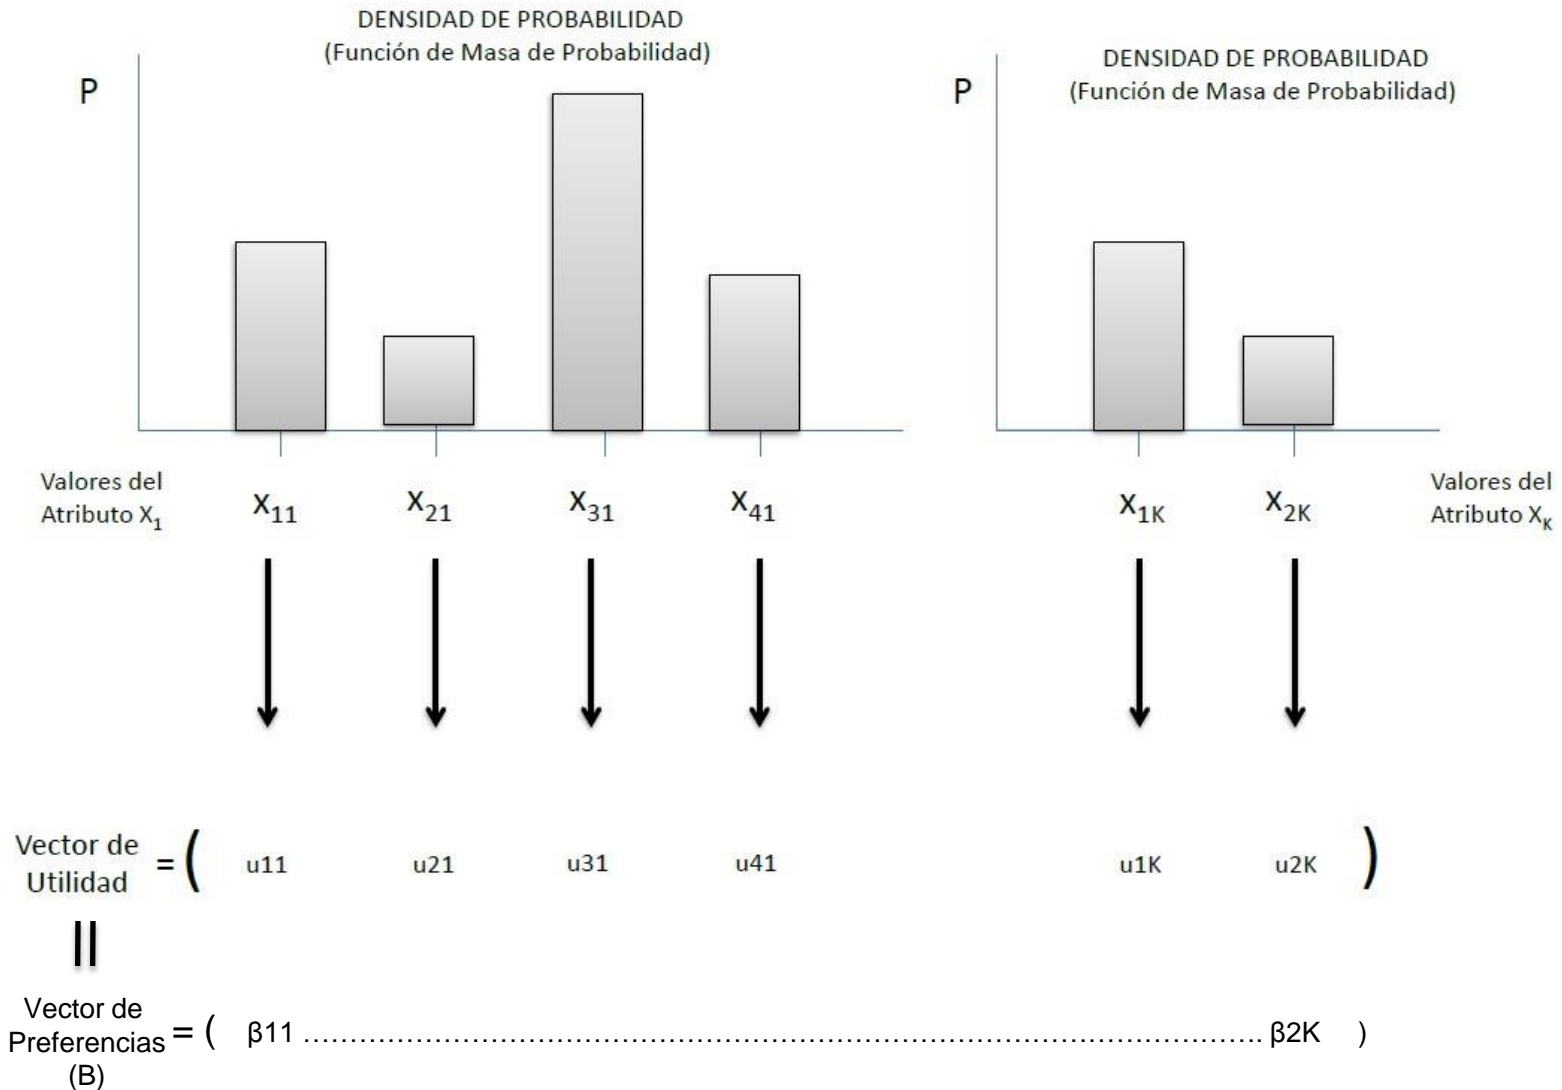
Conjunto de datos (dataset)


Instancias

Atributos

Producto	Autor	Fecha	Precio	Género	Interacción
Libro1	Valor	Valor	Valor	Ficción	Comprado
Libro2	Valor	Valor	Valor	Histórico	No comprado
....
LibroN	Valor	Valor	Valor	Novela	Comprado

Aprendizaje de preferencias (III)






Medida del coseno:

$$\text{similarity} = \cos(\theta) = \frac{A \cdot B}{\|A\| \|B\|} = \frac{\sum_{i=1}^n A_i \times B_i}{\sqrt{\sum_{i=1}^n (A_i)^2} \times \sqrt{\sum_{i=1}^n (B_i)^2}}$$

Utilidad = Similaridad(B,X) = SimilaridadCoseno(B,X)



Distancia Euclídea entre dos vectores:

$$\delta(a, b) = \|a - b\| = \sqrt{(a - b)^T (a - b)} = \sqrt{\sum_{i=1}^m (a_i - b_i)^2}$$

$$\text{Utilidad} = \text{Similaridad}(\mathbf{B}, \mathbf{X}) = 1 / (1 + d(\mathbf{B}, \mathbf{X}))$$

- Most CB-recommendation techniques were applied to recommending text documents.
 - Like web pages or newsgroup messages for example.
- Content of items can also be represented as text documents.
 - With textual descriptions of their basic characteristics.
 - Structured: Each item is described by the same set of attributes



Title	Genre	Author	Type	Price	Keywords
The Night of the Gun	Memoir	David Carr	Paperback	29.90	Press and journalism, drug addiction, personal memoirs, New York
The Lace Reader	Fiction, Mystery	Brunonia Barry	Hardcover	49.90	American contemporary fiction, detective, historical
Into the Fire	Romance, Suspense	Suzanne Brockmann	Hardcover	45.90	American fiction, murder, neo-Nazism

- Unstructured: free-text description.

Item representation

Title	Genre	Author	Type	Price	Keywords
The Night of the Gun	Memoir	David Carr	Paperback	29.90	Press and journalism, drug addiction, personal memoirs, New York
The Lace Reader	Fiction, Mystery	Brunonia Barry	Hardcover	49.90	American contemporary fiction, detective, historical
Into the Fire	Romance, Suspense	Suzanne Brockmann	Hardcover	45.90	American fiction, murder, neo-Nazism

User profile

Title	Genre	Author	Type	Price	Keywords
...	Fiction	Brunonia, Barry, Ken Follett	Paperback	25.65	Detective, murder, New York

$keywords(b_j)$
describes Book b_j
with a set of
keywords



Simple approach

- Compute the similarity of an unseen item with the user profile based on the keyword overlap (e.g. using the Dice coefficient)



$$\frac{2 \times |keywords(b_i) \cap keywords(b_j)|}{|keywords(b_i)| + |keywords(b_j)|}$$

- Or use and combine multiple metrics

Paso 1: Aprender el perfil del usuario (conjunto de términos clave preferentes)

Paso 2: Calcular la similaridad con el item (Tanimoto, Dice, etc.)

Paso 3: Clasificar el item en base a una función umbral:

Relevante? = Si, si Similaridad \Rightarrow R

Relevante? = No, en caso contrario

Term Frequency – Inverse Document Frequency

- Simple keyword representation has its problems
 - in particular when automatically extracted as
 - not every word has similar importance
 - longer documents have a higher chance to have an overlap with the user profile
- Standard measure: TF-IDF
 - Encodes text documents in multi-dimensional Euclidian space
 - weighted term vector
 - TF: Measures, how often a term appears (density in a document)
 - assuming that important terms appear more often
 - normalization has to be done in order to take document length into account
 - IDF: Aims to reduce the weight of terms that appear in all documents

- **Given a keyword i and a document j**
- $TF(i, j)$
 - term frequency of keyword i in document j
- $IDF(i)$
 - inverse document frequency calculated as $IDF(i) = \log \frac{N}{n(i)}$
 - N : number of all recommendable documents
 - $n(i)$: number of documents from N in which keyword i appears
- $TF - IDF$
 - is calculated as: $TF-IDF(i, j) = TF(i, j) * IDF(i)$

- **Term frequency:**

- Each document is a **count vector** in $\mathbb{N}^{|v|}$

	Antony and Cleopatra	Julius Caesar	The Tempest	Hamlet	Othello	Macbeth
Antony	157	73	0	0	0	0
Brutus	4	157	0	1	0	0
Caesar	232	227	0	2	1	1
Calpurnia	0	10	0	0	0	0
Cleopatra	57	0	0	0	0	0
mercy	1.51	0	3	5	5	1
worser	1.37	0	1	1	1	0

Vector v with dimension $|v| = 7$

Combined TF-IDF weights

- Each document is now represented by a real-valued vector of *TF-IDF* weights $\in \mathbb{R}^{|V|}$

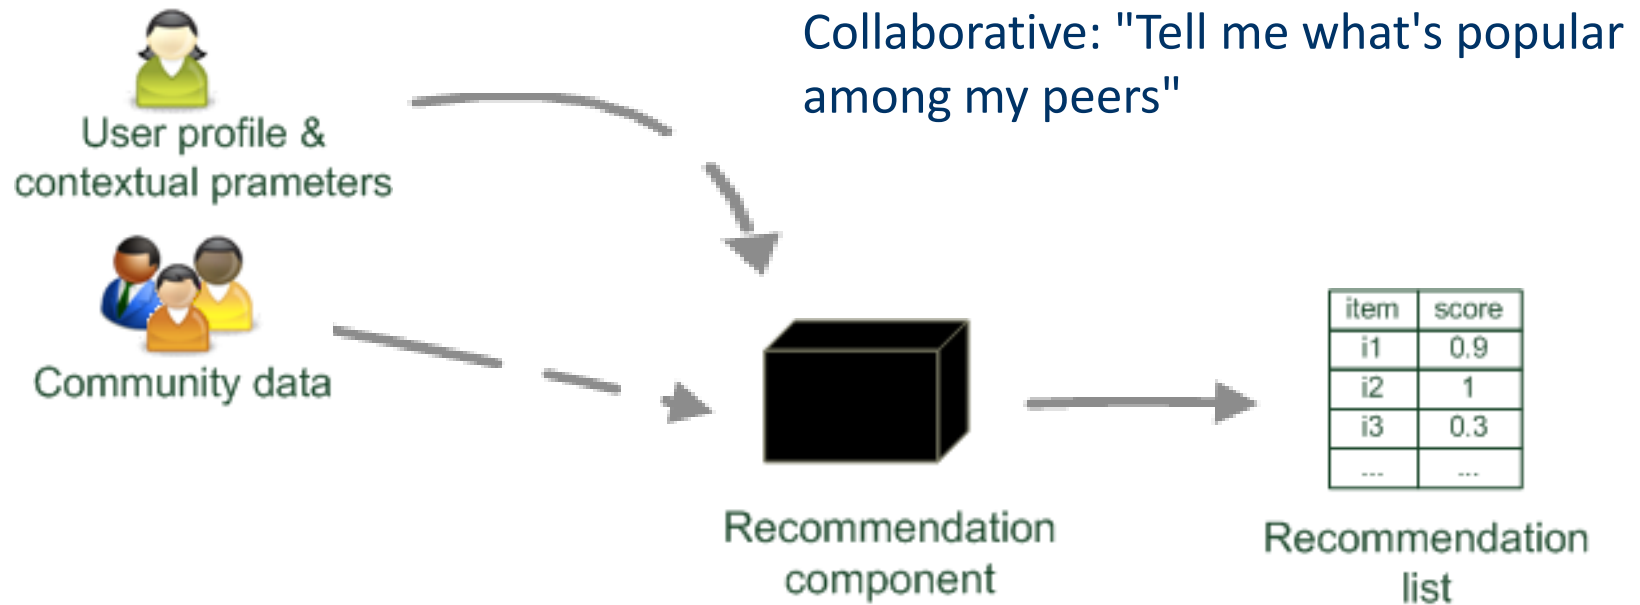
	Antony and Cleopatra	Julius Caesar	The Tempest	Hamlet	Othello	Macbeth
Antony	157	73	0	0	0	0
Brutus	4					
Caesar	232					
Calpurnia	0					
Cleopatra	57					
mercy	1.51					
worser	1.37					

	Antony and Cleopatra	Julius Caesar	The Tempest	Hamlet	Othello	Macbeth
Antony	5.25	3.18	0	0	0	0.35
Brutus	1.21	6.1	0	1	0	0
Caesar	8.59	2.54	0	1.51	0.25	0
Calpurnia	0	1.54	0	0	0	0
Cleopatra	2.85	0	0	0	0	0
mercy	1.51	0	1.9	0.12	5.25	0.88
worser	1.37	0	0.11	4.15	0.25	1.95

Example taken from <http://informationretrieval.org>

Limitaciones:

1. **Para el usuario:** Poca originalidad de las recomendaciones. Las recomendaciones suelen ser productos muy similares a los ya consumidos por el usuario.
2. **Para el ingeniero/científico:** Necesidad de conocer en detalle el dominio de la aplicación: productos, atributos y valores
3. **Para ambos:** Los valores de los atributos no aportan información acerca de la calidad del producto. Estimar la utilidad de un producto a través de sus atributos no garantiza la satisfacción o utilidad de la recomendación.



- The most prominent approach to generate recommendations
 - used by large, commercial e-commerce sites
 - well-understood, various algorithms and variations exist
 - applicable in many domains (book, movies, DVDs, ..)
- Approach
 - use the "wisdom of the crowd" to recommend items
- Basic assumption and idea
 - Users give ratings to catalog items (implicitly or explicitly)
 - Customers who had similar tastes in the past, will have similar tastes in the future
 - Users that rate items similarly have the same tastes



- Input
 - Only a matrix of given user–item ratings
- Output types
 - A (numerical) prediction indicating to what degree the current user will like or dislike a certain item
 - A top-N list of recommended items

		<i>Items</i>					
		<i>1</i>	<i>2</i>	...	<i>i</i>	...	<i>m</i>
<i>Users</i>	<i>1</i>	5	3		1	2	
	<i>2</i>		2				4
	:			5			
	<i>u</i>	3	4		2	1	
	:					4	
	<i>n</i>			3	2		
<i>a</i>		3	5		?	1	

- **The basic technique**

- Given an "active user" (Alice) and an item i not yet seen by Alice
 - find a set of users (peers/nearest neighbors) who liked the same items as Alice in the past **and** who have rated item i
 - use, e.g. the average of their ratings to predict, if Alice will like item i
 - do this for all items Alice has not seen and recommend the best-rated

- **Basic assumption and idea**

- If users had similar tastes in the past they will have similar tastes in the future
- User preferences remain stable and consistent over time

- Example
 - A database of ratings of the current user, Alice, and some other users is given:
 - Ejercicio: ¿Predicción del rating de Alice sobre el Item5?

	Item1	Item2	Item3	Item4	Item5
Alice	5	3	4	4	?
User1	3	1	2	3	3
User2	4	3	4	3	5
User3	3	3	1	5	4
User4	1	5	5	2	1

- Some first questions
 - How do we measure similarity?
 - How many neighbors should we consider?
 - How do we generate a prediction from the neighbors' ratings?



	Item1	Item2	Item3	Item4	Item5
Alice	5	3	4	4	?
User1	3	1	2	3	3
User2	4	3	4	3	5
User3	3	3	1	5	4
User4	1	5	5	2	1

- **A popular similarity measure in user-based CF: Pearson correlation**

a, b : users

$r_{a,p}$: rating of user a for item p

P : set of items, rated both by a and b

- Possible similarity values between -1 and 1

$$\text{sim}(a, b) = \frac{\sum_{p \in P} (r_{a,p} - \bar{r}_a)(r_{b,p} - \bar{r}_b)}{\sqrt{\sum_{p \in P} (r_{a,p} - \bar{r}_a)^2} \sqrt{\sum_{p \in P} (r_{b,p} - \bar{r}_b)^2}}$$

- Ejercicio: Calcula la similaridad utilizando la correlación de Pearson y los datos anteriores

- **A popular similarity measure in user-based CF: Pearson correlation**

a, b : users

$r_{a,p}$: rating of user a for item p

P : set of items, rated both by a and b

- Possible similarity values between -1 and 1

	Item1	Item2	Item3	Item4	Item5
Alice	5	3	4	4	?
User1	3	1	2	3	3
User2	4	3	4	3	5
User3	3	3	1	5	4
User4	1	5	5	2	1




sim = 0,85

sim = 0,70

sim = 0

sim = -0,79

- Ejercicio: ¿Otras formas de calcular la similitud?
- Compara los resultados con la similitud obtenida por correlación de Pearson



- Media ponderada para predecir la valoración:

$$p_{a,i} = \frac{\sum_{j \in K} r_{a,j} w_{i,j}}{\sum_{j \in K} |w_{i,j}|}$$

- $w_{ij} = \text{similaridad}(i,j)$

- A common prediction function:

$$pred(a, p) = \bar{r}_a + \frac{\sum_{b \in N} sim(a, b) * (r_{b,p} - \bar{r}_b)}{\sum_{b \in N} sim(a, b)}$$



- Calculate, whether the neighbors' ratings for the unseen item i are higher or lower than their average
- Combine the rating differences – use the similarity with a as a weight
- Add/subtract the neighbors' bias from the active user's average and use this as a prediction

- Not all neighbor ratings might be equally "valuable"
 - Agreement on commonly liked items is not so informative as agreement on controversial items
 - **Possible solution:** Give more weight to items that have a higher variance
- Value of number of co-rated items
 - Use "significance weighting", by e.g., linearly reducing the weight when the number of co-rated items is low
- Case amplification
 - Intuition: Give more weight to "very similar" neighbors, i.e., where the similarity value is close to 1.
- Neighborhood selection
 - Use similarity threshold or fixed number of neighbors

Recomendación colaborativa: item-based

- Basic idea:
 - Use the similarity between items (and not users) to make predictions
- Example:
 - Look for items that are similar to Item5
 - Take Alice's ratings for these items to predict the rating for Item5

	Item1	Item2	Item3	Item4	Item5
Alice	5	3	4	4	?
User1	3	1	2	3	3
User2	4	3	4	3	5
User3	3	3	1	5	4
User4	1	5	5	2	1

- Produces better results in item-to-item filtering
- Ratings are seen as vector in n-dimensional space
- Similarity is calculated based on the angle between the vectors

$$\text{sim}(\vec{a}, \vec{b}) = \frac{\vec{a} \cdot \vec{b}}{|\vec{a}| * |\vec{b}|}$$




- **Adjusted cosine similarity**
 - take average user ratings into account, transform the original ratings
 - U : set of users who have rated both items a and b

- A common prediction function:

$$pred(u, p) = \frac{\sum_{i \in ratedItem(u)} sim(i, p) * r_{u,i}}{\sum_{i \in ratedItem(u)} sim(i, p)}$$



- Neighborhood size is typically also limited to a specific size
- Not all neighbors are taken into account for the prediction
- An analysis of the MovieLens dataset indicates that "in most real-world situations, a neighborhood of 20 to 50 neighbors seems reasonable" (Herlocker et al. 2002)



1. Validación cruzada con k subgrupos
2. Validación cruzada con 2 subgrupos
3. Validación dejando una instancia fuera (Leave one out)

MEAN SQUARED ERROR
(Error cuadrático medio)

$$\text{MSE} = \frac{1}{n} \sum_{i=1}^n \left(x_i - a \right)^2$$

MEAN ABSOLUTE ERROR
(Error absoluto medio)

$$\text{MAE} = \frac{1}{n} \sum_{i=1}^n |f_i - y_i| = \frac{1}{n} \sum_{i=1}^n |e_i|.$$

Problemas: la calidad de los datos


Training Set			
Id	Status	Age	Class
1	Single	20	Bad
2	Single	30	Good
3	Single	50	Bad
4	Single	60	Good
5	Married	20	Good
6	Married	30	Good
7	Married	40	Good
8	Married	50	Good
9	Divorced	40	Bad
10	Divorced	60	Good

Testing Set			
11	Single	40	(Bad)
12	Married	60	(Good)
13	Divorced	20	(Bad)
14	Divorced	30	(Bad)
15	Divorced	50	(Good)


1. Exactitud de los datos?

2. Imparcialidad del Testing Set Respecto al Training Set?

3. Completitud de los datos



Interpolación:
Situación deseable



Extrapolación:
Problema complicado

Sistemas predictivos: Sistemas de Recomendación



Eduardo M. Sánchez Vila
eduardo.sanchez.vila@usc.es

CITIUS

Grupo de Sistemas Inteligentes
Universidad de Santiago de Compostela